



JOHNS HOPKINS
BLOOMBERG SCHOOL
of PUBLIC HEALTH


Lecture 7

Classification And Regression Trees (CART)
Bagging and Random Forests

Lecture 6 Review:

► CART: Classification And Regression Trees

- Goal is to generate predictions for linear/continuous or categorical (binary or many than 2 categories) responses
- Algorithm goal: minimize a measure of prediction error
- Breaks the predictor X space into non-overlapping sub-spaces, calls these R_m
- Build an initial large tree then prune the tree such that the final tree has smallest prediction error



- Exploring this within NMES, predicting $\log(\text{totalexp}+1)$ and big expenditure

Regression Trees

- Approximates $E(Y|X)$ via a step function!

$$\underbrace{f(X)}_{E(Y|X)} = \sum_{m=1}^M c_m \underbrace{I(X \in R_m)}_{\text{step function}}.$$

- $\underbrace{R_m}$ are selected to minimize

$$\sum_{i=1}^n \left\{ y_i - \sum_{m=1}^M c_m I(X_i \in R_m) \right\}^2$$

- c_m are estimated via the mean Y in R_m

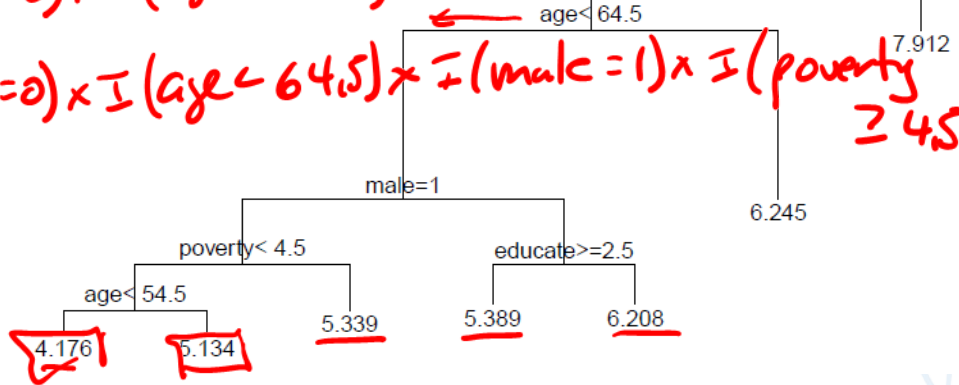
$$\hat{c}_m = \text{ave}(y_i | \underbrace{x_i \in R_m})$$



Example: Predict $\log(\text{expenditure} + 1)$ within NMES

A CART can be translated into a regression model! You do....

$$\begin{aligned} \hat{E}(e | X) = & 7.912 I(\text{mscd}=1) \\ & + 6.245 I(\text{mscd}=0) \times I(\text{age} \geq 64.5) \\ & + 6.208 I(\text{mscd}=0) \times I(\text{age} < 64.5) \times I(\text{male}=0) \times I(\text{educate} < 2.5) \\ & + 5.389 I(\text{mscd}=0) \times I(\text{age} < 64.5) \times I(\text{male}=0) \times I(\text{educate} \geq 2.5) \\ & + 5.339 I(\text{mscd}=0) \times I(\text{age} < 64.5) \times I(\text{male}=1) \times I(\text{poverty} \geq 4.5) \end{aligned}$$



Regression trees: comparing to a parametric model

```
# Fit a parametric model, using the training data
```

```
model=lm(data=dat.train,e~ns(age,2)*male*mscd + as.factor(poverty)*as.factor(educate))
```


```
# Get predictions, residuals based on test/validation sample
```

```
model.yhat=predict(model,newdata=dat.test)
```

```
res.model.test=dat.test$e-model.yhat
```

```
# Compute the MSE for the parametric model
```

```
mse.model.test=sum(res.model.test^2)/length(res.model.test)
```



MSE comparison:

CART: 5.806; Linear Model: 5.68

CART: General comments

Some comments about CART to read further about in THF Chapter 9.2.4:

✂ interactions at the core

- ✂ • must produce rectangles in X space
- ✂ • specific tree, but not necessarily predictions, are unstable to perturbations in X - makes interpretation of tree unreliable
- ✂ • poorly represents smooth functional relationships
- ✂ has natural extensions to the GLM family
- ✂ handles missing data reasonably well through surrogate variables
- ✂ • tends to favor variable selection for factors with lots of levels



CART approach to missing data: surrogate variables

```
## Node number 1: 6151 observations,      complexity param=0.07534598
##   mean=5.908773, MSE=6.708663
##   left son=2 (5463 obs) right son=3 (688 obs)
##   Primary splits:
##       mscd   splits as LR,           improve=0.075345980, (0 missing)
##       age    < 64.5 to the left,      improve=0.047221950, (0 missing)
##       male   splits as RL,           improve=0.008848244, (0 missing)
##       married splits as LRLLL,        improve=0.007339919, (0 missing)
##       beltuse < 2.5 to the left,      improve=0.005627789, (0 missing)
##   Surrogate splits:
##       age < 93.5 to the left, agree=0.888, adj=0.003, (0 split)
##
## Node number 2: 5463 observations,      complexity param=0.02686116
##   mean=5.656467, MSE=6.668967
##   left son=4 (3446 obs) right son=5 (2017 obs)
##   Primary splits:
##       age    < 64.5 to the left,      improve=0.030424030, (0 missing)
##       male   splits as RL,           improve=0.013146610, (0 missing)
##       beltuse < 2.5 to the left,      improve=0.006562917, (0 missing)
##       educate < 2.5 to the right,    improve=0.006149474, (0 missing)
##       married splits as LRLLL,        improve=0.004817598, (0 missing)
##   Surrogate splits:
##       married splits as LRLLL,        agree=0.715, adj=0.229, (0 split)
##       educate < 3.5 to the left,      agree=0.658, adj=0.074, (0 split)
##   ...
```

Classification trees

- ▶ Same procedure
- ▶ Different goodness of fit criteria

For classification trees, define

$$\hat{p}_{mk} = \frac{1}{N_m} \sum_{x_i \in R_m} I(y_i = k)$$

$y_i = \text{Binary}$
 $k = 0, 1$

$R_m: \hat{p}_{m0}, \hat{p}_{m1}$

Observations in node m are classified in category k based on the category k with highest \hat{p}_{mk} .

Different measures of node impurity $Q_m(T)$ include the misclassification error, Gini index and cross-entropy or deviance. When we are constructing a classification tree for a binary response where $p = \text{Pr}(Y = 1)$, these measures are:

* Misclassification error: $1 - \max(p, 1 - p)$

* Gini index: $2p(1 - p)$

* Cross entropy or deviance: $-p \log(p) - (1 - p) \log(1 - p)$


$R_m: \hat{p}_{m0} = .7, \hat{p}_{m1} = .3$
classification $R_m = 0$
misclassification
error = .3

The Gini index or cross-entropy impurity measures are often used to construct the tree where the misclassification error is used for tree pruning.

Example: Predict a big expenditure


```
set.seed(123454321)
dat1.train=dat1[train<-sample(1:nrow(dat),floor(nrow(dat)/2)),]
dat1.test=dat1[-train,]
tree0=rpart(big~.,data=dat1.train,method="class",control=rpart.control(minsize=20,cp=.001))
par(mfrow=c(2,1),mar=c(5,5,5,1))
plotcp(tree0)
```

*big = { 1 7,000
0 0/w*



Example: Predict a big expenditure

```
tree=rpart(big~.,data=dat1.train,method="class",control=rpart.control(minsize=20,cp=.02))  
plotcp(tree)
```




	CP	nsplit	rel error	xerror	xstd
##	1	0.1555751	0	1.0000000	0.01667776
##	2	0.0200000	1	0.8444249	0.01600726

Variable importance

mscd

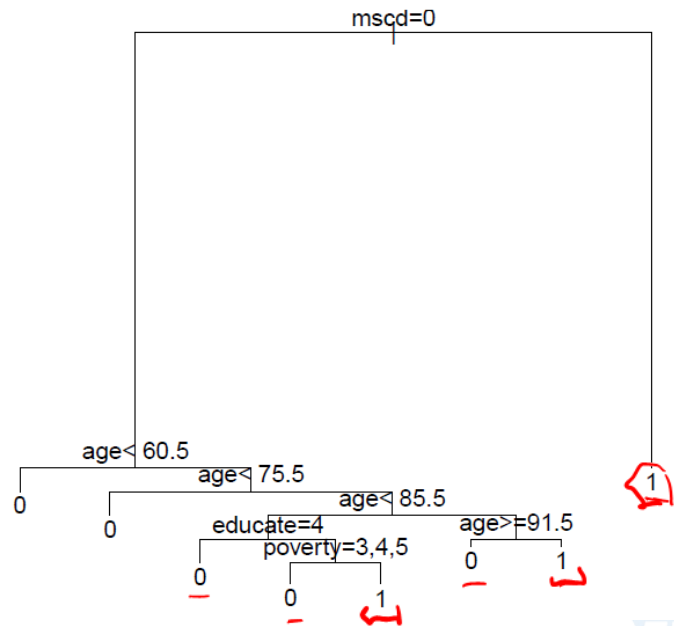
100



Example: Predict a big expenditure

► Change the CART criteria

```
treetest0=rpart(big~.,data=dat1.train,method="class",control=rpart.control(minsize=5,cp=.001))  
bestcp <- treetest0$cptable[which.min(treetest0$cptable[, "xerror"]), "CP"]  
tree.pruned <- prune(treetest0, cp = bestcp)  
plot(tree.pruned);text(tree.pruned,pretty=3)
```



Example: Predict a big expenditure

- Compare to a parametric model

```
# comparison to logistic model
model=glm(data=dat1.train,big~ns(age,2)*male*mscd
          + as.factor(poverty)*as.factor(educate),
          family=binomial())

# Get the predicted values from the logistic regression model
model.yhat=predict.glm(model,newdata=dat1.test,type="response")
# The ROCR package requires that you first create a "prediction"
# object using the prediction function,
# this function takes the predicted probabilities + true values
pred.model.test=prediction(model.yhat,dat1.test$big)
# The performance function will compute several measures of
# performance for the classification scheme
# here we are selecting true positive rate, false positive rate
perf.model.test=performance(pred.model.test,"tpr","fpr")
# Use the performance object and ask for AUC
auc.model.test=performance(pred.model.test,"auc")
```

Example: Predict a big expenditure

```
# Compare the performance of the parametric model  
# to the classification tree mnsize = 20  
tree.yhat=as.vector(predict(tree,newdata=dat1.test,na.action=na.pass)[,2])  
# Same as before, create the prediction object and  
# compute auc  
pred.tree.test=prediction(tree.yhat,dat1.test$big)  
auc.tree.test=performance(pred.tree.test,"auc")
```

$$\hat{P}_r(Y_i=0), \hat{P}_r(Y_i=1)$$

↓
 \hat{P}_m

```
# Compare the performance of the parametric model  
# to the classification tree mnsize = 5  
tree.yhat2=as.vector(predict(tree.pruned,newdata=dat1.test,na.action=na.pass)[,2])  
# Same as before, create the prediction object and  
# compute auc  
pred.tree.test2=prediction(tree.yhat2,dat1.test$big)  
auc.tree.test2=performance(pred.tree.test2,"auc")
```

Example: Predict a big expenditure

- Compare the AUC values

```
auc.model.test@y.values[[1]]
```

```
## [1] 0.6762746
```

```
auc.tree.test@y.values[[1]]
```

```
## [1] 0.588122
```

```
auc.tree.test2@y.values[[1]]
```

```
## [1] 0.6463928
```

parametric model

→ 1st tree ⇒ mscd

*→ 2nd tree
mscd, age, education*



HEART: Foundations of Statistical Machine Learning

Class 8: November 8th, 2018




Shannon Wongvibulsin
Johns Hopkins School of Medicine
Biomedical Engineering Department
MD/PhD Candidate



Ensemble Methods

- **Bagging (Bootstrap Aggregating)**
- **Random Forest**



Ensemble Methods

- **Bagging (Bootstrap Aggregating)**
- **Random Forest**





Image source:


<https://www.google.com/url?sa=i&source=images&cd=&cad=rja&uact=8&ved=2ahUKEwjMmruX-cLeAhVrmuAKHXETDpQQjRx6BAgBEAU&url=https%3A%2F%2Fcloudblogs.microsoft.com%2Fmicrosoftsecure%2F2018%2F02%2F14%2Fhow-artificial-intelligence-stopped-an-emetet-outbreak%2F&psig=AOvVaw20oWhcs-VEjNXILgXX-cVW&ust=1541702921680011>

Bootstrapping



Bagging

Create B bootstrap samples
by sampling with replacement
from the training sample



n : number of observations
 p : number of predictors

Bagging

Create **B** bootstrap samples
by sampling with replacement
from the training sample

Training Sample
 $n \times p$

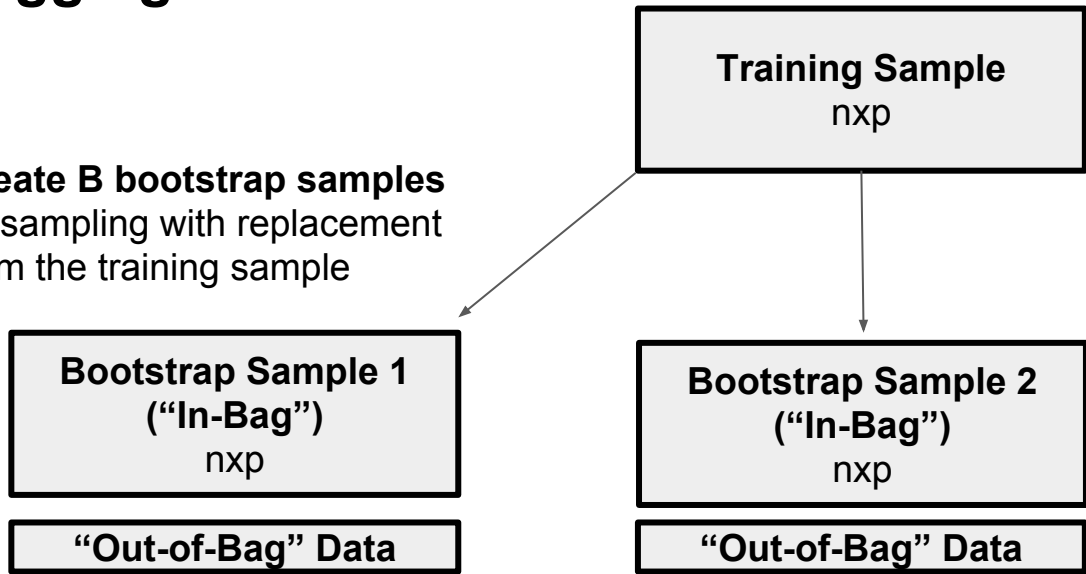
n : number of observations
 p : number of predictors

Bootstrap Sample 1
("In-Bag")
 $n \times p$

"Out-of-Bag" Data

Bagging

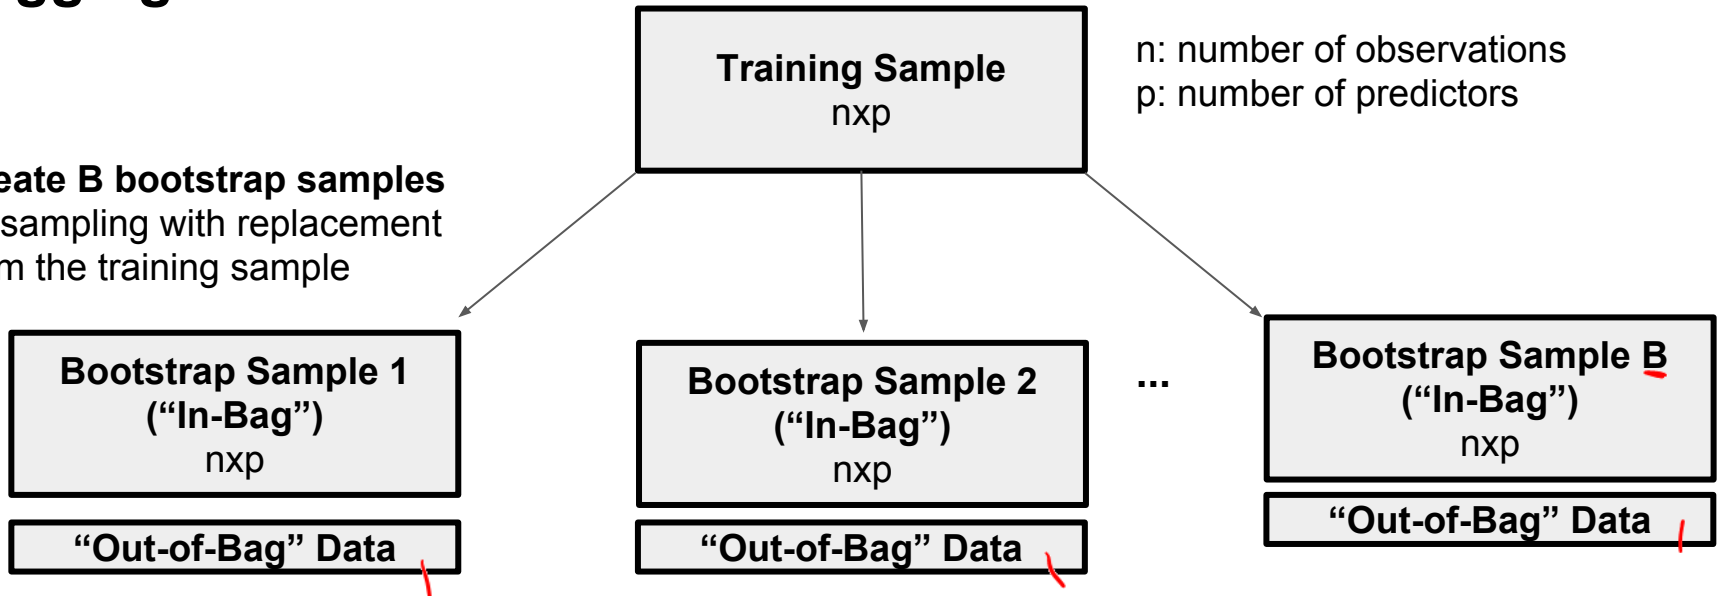
Create **B** bootstrap samples
by sampling with replacement
from the training sample



n : number of observations
 p : number of predictors

Bagging

Create **B** bootstrap samples
by sampling with replacement
from the training sample



Bagging

Create **B** bootstrap samples
by sampling with replacement
from the training sample




n : number of observations
 p : number of predictors



...



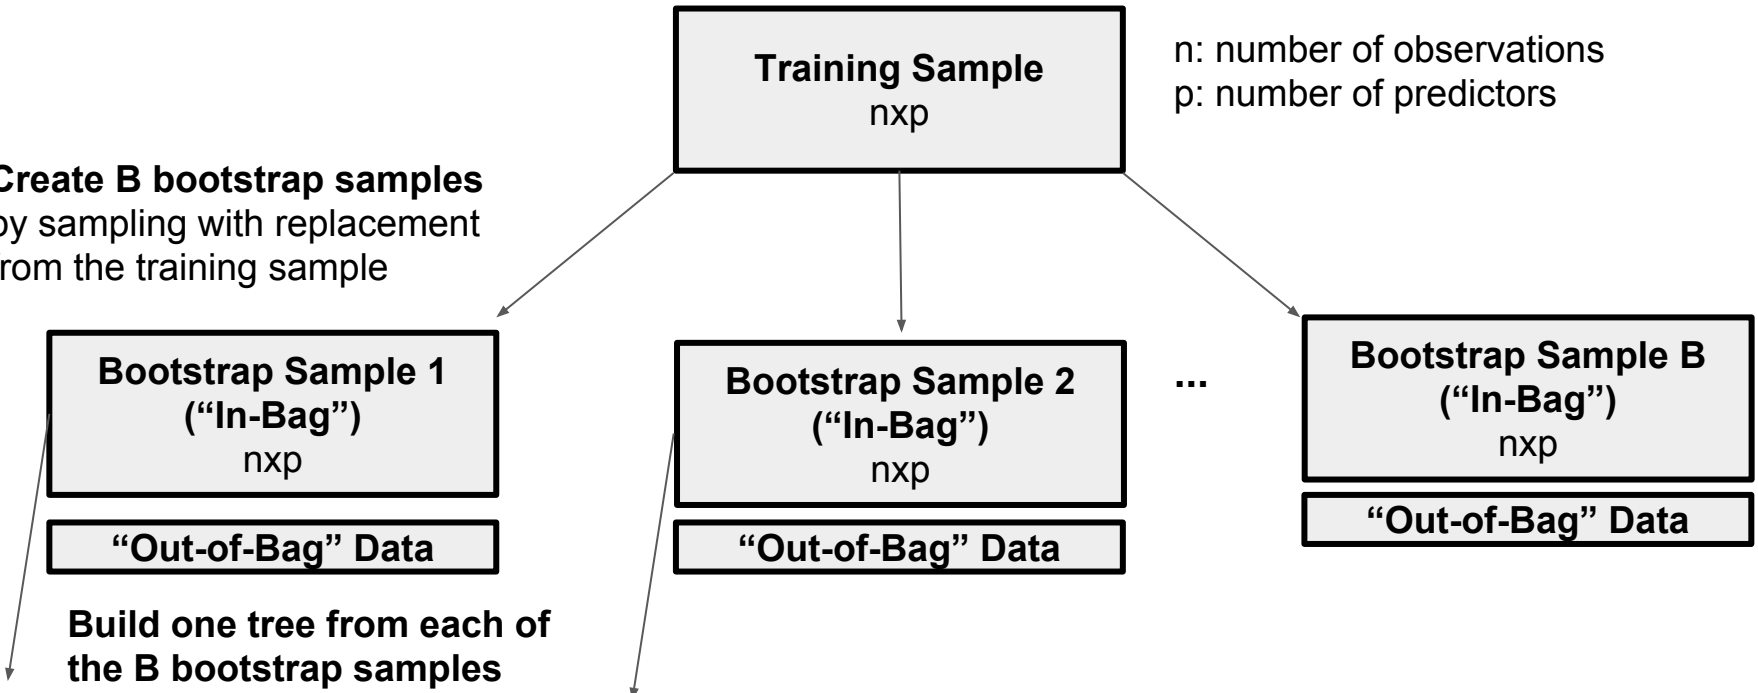


Build one tree from each of
the **B** bootstrap samples



Random Forests

Create **B** bootstrap samples by sampling with replacement from the training sample



Bagging

Create **B** bootstrap samples by sampling with replacement from the training sample






n : number of observations
 p : number of predictors



...

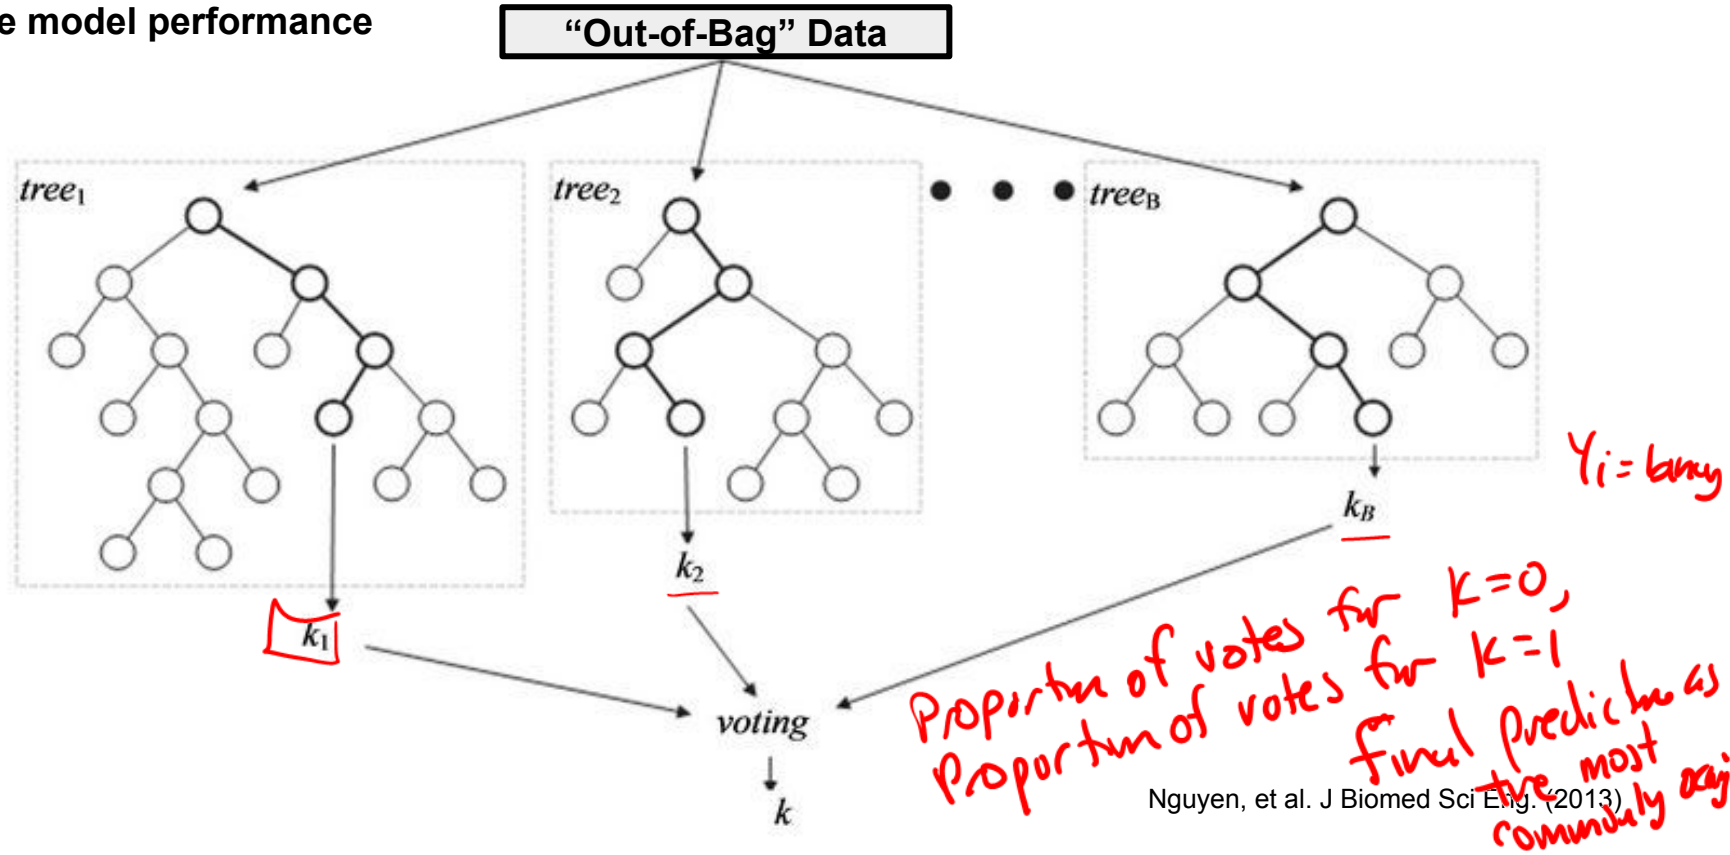


Build one tree from each of the **B** bootstrap samples



Bagging

Use “Out-of-Bag” (OOB) data to estimate model performance



Bagging

Issue with decision trees: **high variance (overfitting)**

- Example:
 - Build decision trees on data split in random different ways
 - Decision trees give different results (high variance)

Addressing the high variance problem with bagging:

1. Build decision trees on B **bootstrap*** samples
2. Average predictions over all decision trees

$$\hat{f}_{avg}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x)$$

Note: no pruning of trees

*Recall: **Bootstrapping**:

- Sample with replacement
- In-bag: $\sim 2/3$ of data
- Out-of-bag: $\sim 1/3$ of data

Potential Issue with Bagging

- **Correlated trees**

- Example:

- One very strong predictor

- All bagged trees will select strong predictor at top of tree
 - All bagged trees will be similar

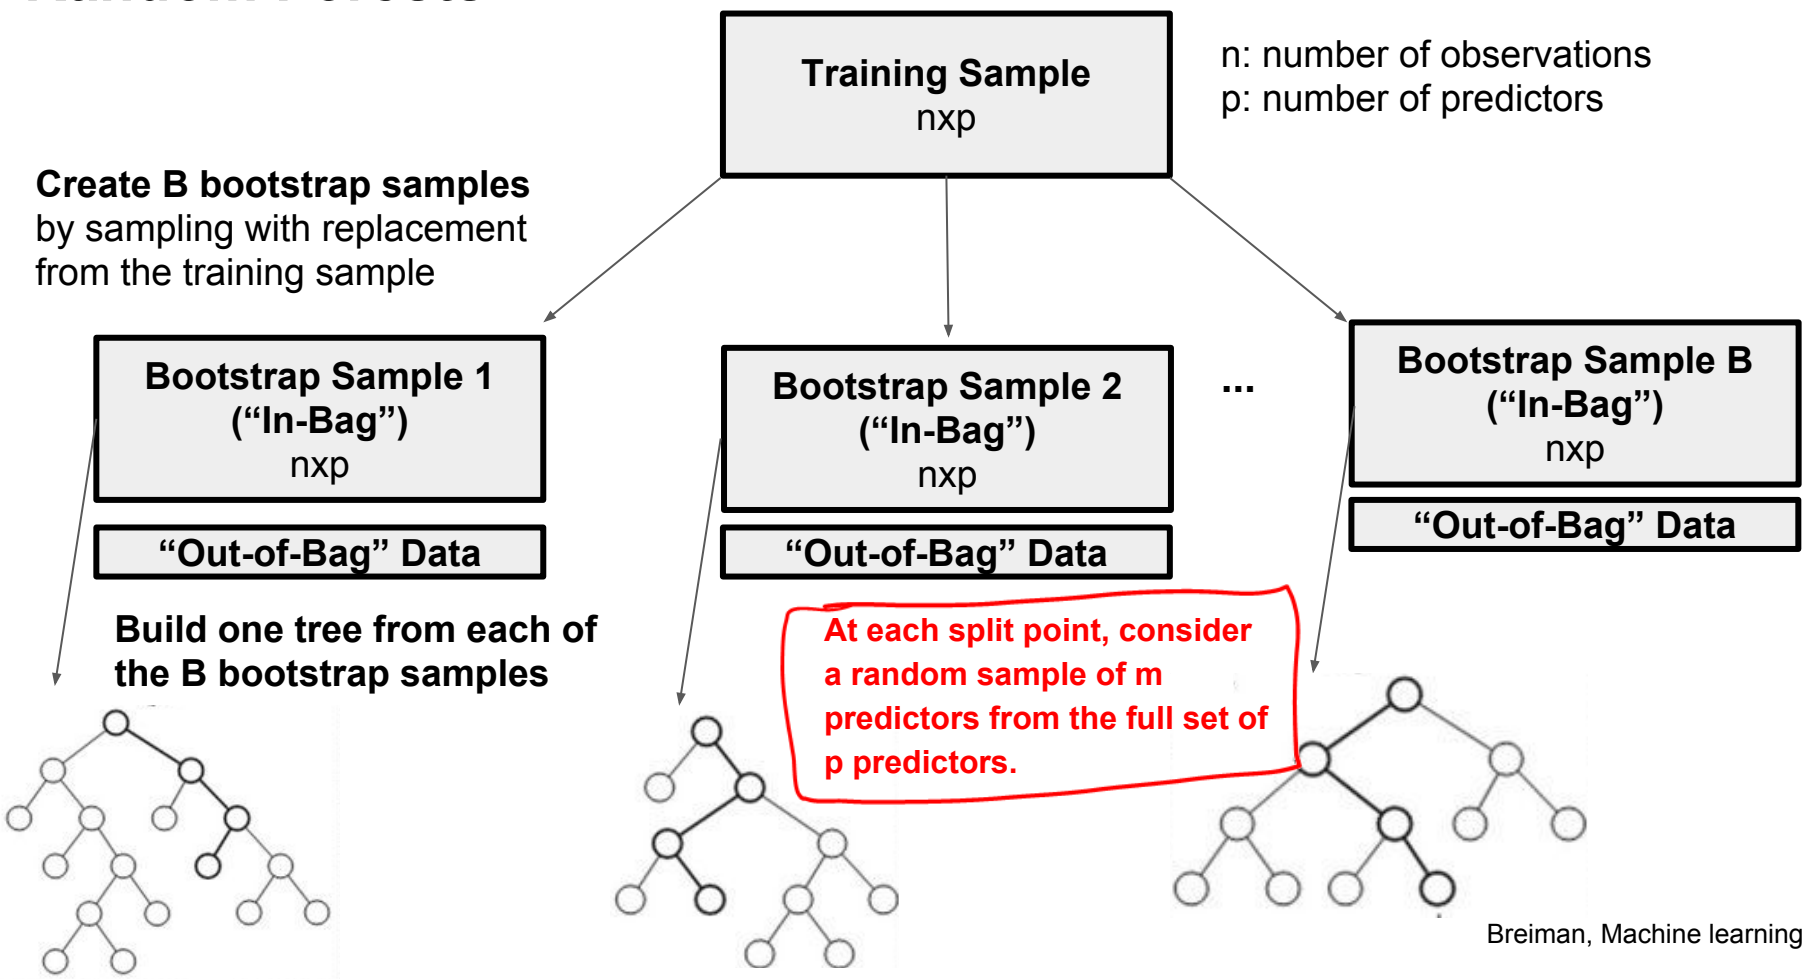
How to decorrelate trees constructed from bootstrap samples?

Random Forest

Random Forests

Create **B** bootstrap samples by sampling with replacement from the training sample

n: number of observations
p: number of predictors

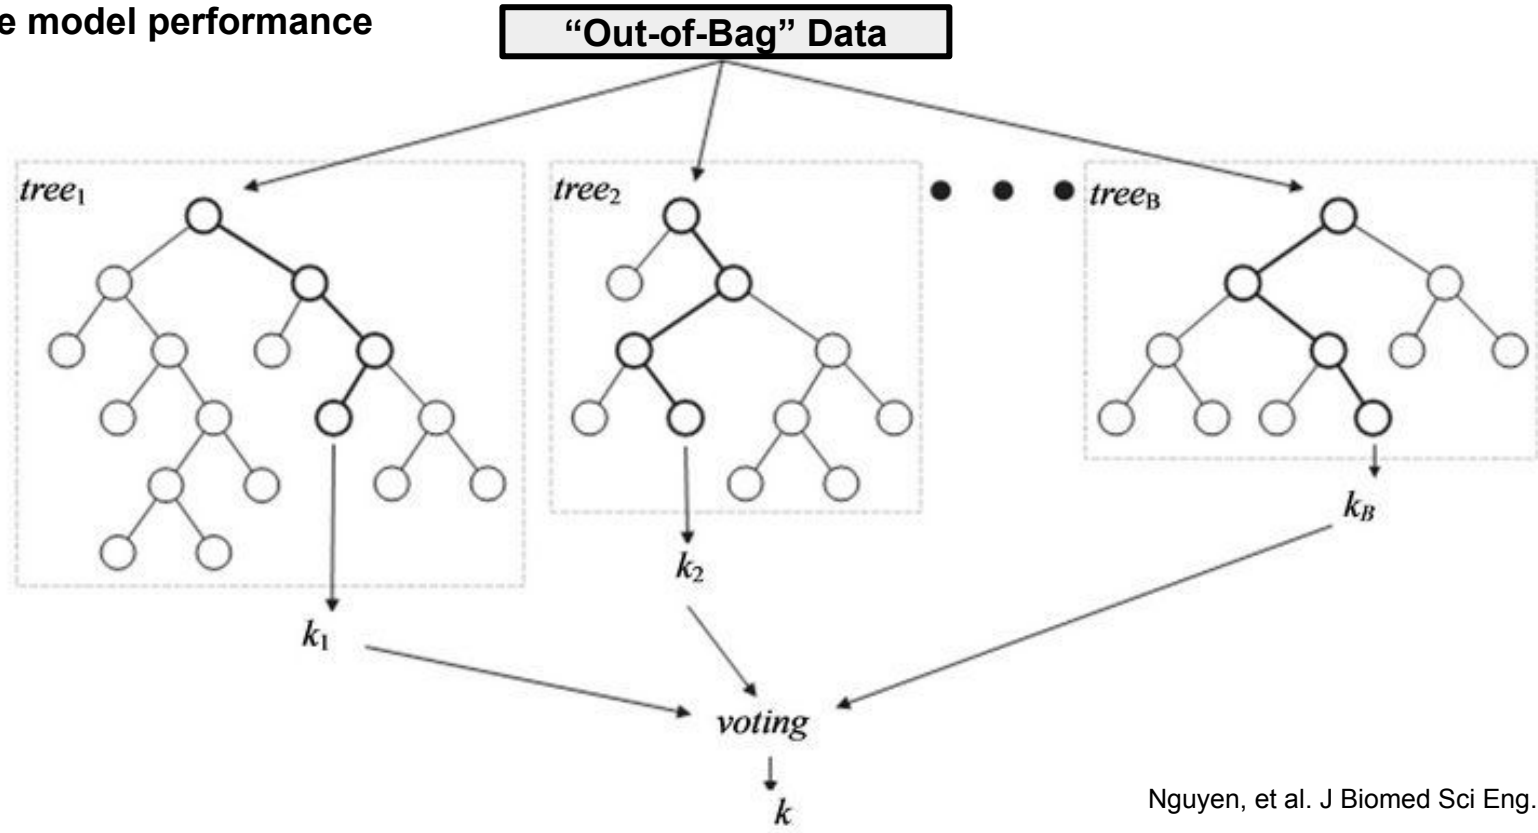


Random Forest

- 1. Build decision trees on B **bootstrap** samples
 - **When a split is considered, a random sample of m predictors is chosen as split candidates from the full set of p predictors.**
 - Note:
 - $m \leq p$
 - If $m = p$, bagging
- 2. Average predictions over all decision trees

Random Forests

Use “Out-of-Bag” (OOB) data to estimate model performance



Random Forest Algorithm, HTF text

Algorithm 15.1 *Random Forest for Regression or Classification.*

1. For $b = 1$ to B :
 - (a) Draw a bootstrap sample \mathbf{Z}^* of size N from the training data.
 - (b) Grow a random-forest tree T_b to the bootstrapped data, by recursively repeating the following steps for each terminal node of the tree, until the minimum node size n_{min} is reached.
 - i. Select m variables at random from the p variables.
 - ii. Pick the best variable/split-point among the m .
 - iii. Split the node into two daughter nodes.
2. Output the ensemble of trees $\{T_b\}_1^B$.



To make a prediction at a new point x :

Regression: $\hat{f}_{\text{rf}}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$.

Classification: Let $\hat{C}_b(x)$ be the class prediction of the b th random-forest tree. Then $\hat{C}_{\text{rf}}^B(x) = \text{majority vote } \{\hat{C}_b(x)\}_1^B$.



Random Forest Extras

- ▶ Missing values in the training data:
 - ▶ Random forests do not like missing values in the training data
 - ▶ First impute missing data:
 - Mean / mode replacement
 - Imputation via proximity; see these youtube videos which do a good job of giving the overview of how the imputation works
 - ▶  https://www.youtube.com/watch?v=J4Wdy0Wc_xQ
 - ▶  https://www.youtube.com/watch?v=nyxTdL_4Q-Q
- ▶ Missing values for testing/validation data:
 - ▶ Surrogate variables are used.
- ▶ Given the internal cross-validation that occurs, do we need to separate data into training and test/validation?
 - ▶ Not really! So long as you evaluate the utility of the random forest using the out-of-bag predictions/error!
- ▶ In PS2, you will be using a training and test/validation because for learning we will ask you to compare the utility of a parametric model, a single CART and a random forest. But for applications where you will use the random forest, you don't need to separate the data.



Random Forest Extras

- ▶ Parameters that we control
 - ▶ Number of variables considered at each split, m
 - Classification tree: floor square-root p
 - Regression tree: floor $p/3$
 - ▶ Number of trees
 - ▶ Recommendation: To find m : set number of trees large (e.g. 500), identify minimum out-of-bag error for $m = 1, 2, \dots$, beyond default. After finding m : check to see if your forest is sensitive to number of trees by plotting MSE or out-of-bag error as a function of number of trees.

* minimum node size = 5

- ▶ Out of bag samples

For each observation $z_i = (x_i, y_i)$, construct its random forest predictor by averaging only those trees corresponding to bootstrap samples in which z_i did not appear.

- ▶ Out of bag error -> corresponds to a n-fold cross-validation

- ▶ Variable importance

- ▶ Ranks each variable by summing up (over all trees) "improvement" in prediction error when variable is included